



# 程序的性能优化

王晶

jwang@ruc.edu.cn, 信息楼124

2025年12月



# 提纲

- 概述
- 常用的优化技术
  - 代码移动/预计算
  - 减小强度
  - 共享相同子表达式
  - 示例：冒泡排序
- 优化的阻碍
  - 函数调用
  - 内存别名
- 指令级并行
- 处理条件判断



# 关于性能

- *There's more to performance than asymptotic complexity*  
(复杂度)
- Constant factors matter too!
  - Easily see 10:1 performance range depending on **how code is written**
  - Must optimize at multiple levels:
    - 算法、数据类型、函数、循环
- Must understand system to optimize performance
  - How programs are compiled and executed, 编译
  - How modern processors + memory systems operate, 架构+存储
  - How to measure program performance and identify bottlenecks, 瓶颈分析
  - How to improve performance without destroying code modularity and generality, 模块化和通用性



# 优化编译器

- Provide efficient mapping of program to machine
  - register allocation寄存器分配
  - code selection and ordering (scheduling)代码选择和调度
  - dead code elimination消除无效代码
  - eliminating minor inefficiencies消除低效代码
- Don't (usually) improve asymptotic efficiency
  - up to programmer to select best overall algorithm
  - big-O savings are (often) more important than constant factors
    - but constant factors also matter
- Have difficulty overcoming “optimization blockers”
  - potential memory aliasing
  - potential procedure side-effects



# 提纲

- 概述
- 常用的优化技术
  - 代码移动/预计算
  - 减小强度
  - 共享相同子表达式
  - 示例：冒泡排序
- 优化的阻碍
  - 函数调用
  - 内存别名
- 指令级并行
- 处理条件判断



# Generally Useful Optimizations

- Optimizations that you or the compiler should do regardless of processor / compiler
- 代码移动 Code Motion
  - Reduce frequency with which computation performed
    - If it will always produce same result
    - Especially moving code out of loop

```
void set_row(double *a, double *b,  
            long i, long n)  
{  
    long j;  
    for (j = 0; j < n; j++)  
        a[n*i+j] = b[j];  
}
```



```
long j;  
int ni = n*i;  
for (j = 0; j < n; j++)  
    a[ni+j] = b[j];
```



# Compiler-Generated Code Motion (-O1)

```
void set_row(double *a, double *b,
             long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```

```
long j;
long ni = n*i;
double *rowp = a+ni;
for (j = 0; j < n; j++)
    *rowp++ = b[j];
```

```
set_row:
    testq    %rcx, %rcx                # Test n
    jle      .L1                      # If 0, goto done
    imulq    %rcx, %rdx                # ni = n*i
    leaq     (%rdi,%rdx,8), %rdx        # rowp = A + ni*8
    movl     $0, %eax                  # j = 0
.L3:
    movsd    (%rsi,%rax,8), %xmm0      # t = b[j]
    movsd    %xmm0, (%rdx,%rax,8)      # M[A+ni*8 + j*8] = t
    addq     $1, %rax                  # j++
    cmpq     %rcx, %rax                # j:n
    jne      .L3                      # if !=, goto loop
.L1:
    rep ; ret                          # done:
```



# Reduction in Strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide
  - $16 * x \quad \rightarrow \quad x \ll 4$
  - Utility machine dependent
  - Depends on cost of multiply or divide instruction
    - On Intel Nehalem, integer multiply requires 3 CPU cycles
- Recognize sequence of products

```
for (i = 0; i < n; i++) {  
    int ni = n*i;  
    for (j = 0; j < n; j++)  
        a[ni + j] = b[j];  
}
```



```
int ni = 0;  
for (i = 0; i < n; i++) {  
    for (j = 0; j < n; j++)  
        a[ni + j] = b[j];  
    ni += n;  
}
```





# Share Common Subexpressions

- Reuse portions of expressions
- GCC will do this with `-O1`

```
/* Sum neighbors of i,j */
up =    val[(i-1)*n + j  ];
down =  val[(i+1)*n + j  ];
left =  val[i*n          + j-1];
right = val[i*n          + j+1];
sum = up + down + left + right;
```

3 multiplications:  $i*n$ ,  $(i-1)*n$ ,  $(i+1)*n$

```
long inj = i*n + j;
up =    val[inj - n];
down =  val[inj + n];
left =  val[inj - 1];
right = val[inj + 1];
sum = up + down + left + right;
```

1 multiplication:  $i*n$

```
leaq    1(%rsi), %rax    # i+1
leaq    -1(%rsi), %r8    # i-1
imulq   %rcx, %rsi      # i*n
imulq   %rcx, %rax      # (i+1)*n
imulq   %rcx, %r8      # (i-1)*n
addq    %rdx, %rsi      # i*n+j
addq    %rdx, %rax      # (i+1)*n+j
addq    %rdx, %r8      # (i-1)*n+j
...
```

```
imulq   %rcx, %rsi      # i*n
addq    %rdx, %rsi      # i*n+j
movq    %rsi, %rax      # i*n+j
subq    %rcx, %rax      # i*n+j-n
leaq    (%rsi,%rcx), %rcx # i*n+j+n
...
```



# Bubblesort

比较遍数

| 数据   | 1    | 2    | 3    | 4    |
|------|------|------|------|------|
| 587  | -632 | -632 | -632 | -632 |
| -632 | 587  | 234  | -34  | -34  |
| 777  | 234  | -34  | 234  | 234  |
| 234  | -34  | 587  | 587  | 587  |
| -34  | 777  | 777  | 777  | 777  |

从小到大排序



# Optimization Example: Bubblesort

- **Bubblesort** program that sorts an array **A** that is allocated in static storage:
  - an element of **A** requires **four bytes** of a byte-addressed machine
  - elements of **A** are numbered **1 through n** (**n** is a variable)
  - **A[j]** is in location **&A+4\*(j-1)**

```
for (i = n-1; i >= 1; i--) {  
    for (j = 1; j <= i; j++)  
        if (A[j] > A[j+1]) {  
            temp = A[j];  
            A[j] = A[j+1];  
            A[j+1] = temp;  
        }  
}
```



# Translated (Pseudo) Code

```
i := n-1
L5:  if i<1 goto L1
    j := 1
L4:  if j>i goto L2
    t1 := j-1
    t2 := 4*t1
    t3 := A[t2]    // A[j]
    t4 := j+1
    t5 := t4-1
    t6 := 4*t5
    t7 := A[t6]    // A[j+1]
    if t3<=t7 goto L3
```

```
for (i = n-1; i >= 1; i--) {
  for (j = 1; j <= i; j++)
    if (A[j] > A[j+1]) {
      temp = A[j];
      A[j] = A[j+1];
      A[j+1] = temp;
    }
}
```

```
t8 := j-1
t9 := 4*t8
temp := A[t9]    // temp:=A[j]
t10 := j+1
t11:= t10-1
t12 := 4*t11
t13 := A[t12]    // A[j+1]
t14 := j-1
t15 := 4*t14
A[t15] := t13    // A[j]:=A[j+1]
t16 := j+1
t17 := t16-1
t18 := 4*t17
A[t18]:=temp    // A[j+1]:=temp
L3: j := j+1
    goto L4
L2: i := i-1
    goto L5
L1:
```

## Instructions

29 in outer loop  
25 in inner loop



# Redundancy in Address Calculation

```
i := n-1
L5:  if i<1 goto L1
    j := 1
L4:  if j>i goto L2
    t1 := j-1
    t2 := 4*t1
    t3 := A[t2] // A[j]
    t4 := j+1
    t5 := t4-1
    t6 := 4*t5
    t7 := A[t6] // A[j+1]
    if t3<=t7 goto L3
```

```
t8 := j-1
t9 := 4*t8
temp := A[t9] // temp:=A[j]
t10 := j+1
t11 := t10-1
t12 := 4*t11
t13 := A[t12] // A[j+1]
t14 := j-1
t15 := 4*t14
A[t15] := t13 // A[j]:=A[j+1]
t16 := j+1
t17 := t16-1
t18 := 4*t17
A[t18] := temp // A[j+1]:=temp
L3:  j := j+1
    goto L4
L2:  i := i-1
    goto L5
L1:
```



# Redundancy Removed

```
i := n-1
L5: if i<1 goto L1
    j := 1
L4: if j>i goto L2
    t1 := j-1
    t2 := 4*t1
    t3 := A[t2]      // A[j]
    t6 := 4*j
    t7 := A[t6]      // A[j+1]
    if t3<=t7 goto L3
```

```
t8 :=j-1
t9 := 4*t8
temp := A[t9]      // temp:=A[j]
t12 := 4*j
t13 := A[t12]      // A[j+1]
A[t9] := t13       // A[j]:=A[j+1]
A[t12] :=temp      // A[j+1]:=temp
L3: j := j+1
    goto L4
L2: i := i-1
    goto L5
L1:
```

Instructions  
20 in outer loop  
16 in inner loop



# More Redundancy

```
i := n-1
L5: if i<1 goto L1
    j := 1
L4: if j>i goto L2
    t1 := j-1
    t2 := 4*t1
    t3 := A[t2]      // A[j]
    t6 := 4*j
    t7 := A[t6]      // A[j+1]
    if t3<=t7 goto L3
```

```
t8 := j-1
t9 := 4*t8
temp := A[t9] // temp:=A[j]
t12 := 4*j
t13 := A[t12] // A[j+1]
A[t9] := t13 // A[j]:=A[j+1]
A[t12] := temp // A[j+1]:=temp
L3: j := j+1
    goto L4
L2: i := i-1
    goto L5
L1:
```



# Redundancy Removed

```
i := n-1
L5: if i<1 goto L1
    j := 1
L4: if j>i goto L2
    t1 := j-1
    t2 := 4*t1
    t3 := A[t2]      // old_A[j]
    t6 := 4*j
    t7 := A[t6]      // A[j+1]
    if t3<=t7 goto L3
L3: j := j+1
    goto L4
L2: i := i-1
    goto L5
L1:
```

**A[t2] := t7** // A[j]:=A[j+1]  
**A[t6] := t3** // A[j+1]:=old\_A[j]

Instructions  
15 in outer loop  
11 in inner loop





# Redundancy in Loops

```
i := n-1  
L5: if i<1 goto L1  
  
j := 1  
L4: if j>i goto L2  
    t1 := j-1  
    t2 := 4*t1  
    t3 := A[t2]    // A[j]  
    t6 := 4*j  
    t7 := A[t6]    // A[j+1]  
    if t3<=t7 goto L3  
    A[t2] := t7  
    A[t6] := t3  
L3: j := j+1  
    goto L4  
L2: i := i-1  
    goto L5  
L1:
```



# Redundancy Eliminated

```
    i := n-1
L5:  if i<1 goto L1
    j := 1
L4:  if j>i goto L2
    t1 := j-1
    t2 := 4*t1
    t3 := A[t2]    // A[j]
    t6 := 4*j
    t7 := A[t6]    // A[j+1]
    if t3<=t7 goto L3
    A[t2] := t7
    A[t6] := t3
L3:  j := j+1
    goto L4
L2:  i := i-1
    goto L5
L1:
```

```
    i := n-1
L5:  if i<1 goto L1
    t2 := 0
    t6 := 4
    t19 := 4*i
L4:  if t6>t19 goto L2
    t3 := A[t2]
    t7 := A[t6]
    if t3<=t7 goto L3
    A[t2] := t7
    A[t6] := t3
L3:  t2 := t2+4
    t6 := t6+4
    goto L4
L2:  i := i-1
    goto L5
L1:
```



# Final Pseudo Code

```
    i := n-1
L5:  if i<1 goto L1
    t2 := 0
    t6 := 4
    t19 := i << 2
L4:  if t6>t19 goto L2
    t3 := A[t2]
    t7 := A[t6]
    if t3<=t7 goto L3
    A[t2] := t7
    A[t6] := t3
L3:  t2 := t2+4
    t6 := t6+4
    goto L4
L2:  i := i-1
    goto L5
L1:
```

## Instruction Count Before Optimizations

**29 in outer loop  
25 in inner loop**

## Instruction Count After Optimizations

**15 in outer loop  
9 in inner loop**

- These were **Machine-Independent Optimizations**.
- Will be followed by **Machine-Dependent Optimizations**, including allocating temporaries to registers, converting to assembly code




# 提纲

- 概述
- 常用的优化技术
  - 代码移动/预计算
  - 减小强度
  - 共享相同子表达式
  - 示例：冒泡排序
- 优化的阻碍
  - 函数调用
  - 内存别名
- 指令级并行
- 处理条件判断



# 优化编译器的能力和局限性

- GCC胜任基本的优化，不适合执行激进的变换
- GCC程序员需要强化代码优化，以简化编译器生成高效代码任务的方式编写程序



## GCC-conscious programming



# 编译器优化的能力和局限性

- 优化行为受到基本的约束
  - 不能改变程序行为
    - 除非程序使用非标准语言功能时
  - 通常会阻止它进行仅影响特殊条件下行为的优化
- 对程序员来说可能很明显的行为可能会被语言和编码风格混淆
  - e.g., Data ranges may be more limited than variable types suggest 例如，数据范围可能比变量类型的建议范围更有限
- 大多数分析仅在程序内执行
  - 在大多数情况下，整个程序分析的成本太高
  - 较新版本的 GCC 在单个文件中执行过程间分析
    - 但是，不是在不同文件中的代码之间
- 大多数分析仅基于静态信息
  - 编译器难以预测运行时输入
- 如有疑问，编译器必须是保守的



# Optimization Blocker #1: Procedure Calls

- Procedure to Convert String to Lower Case


```
void lower(char *s)
{
    size_t i;
    for (i = 0; i < strlen(s); i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

- Extracted from 213 lab submissions, Fall, 1998



# Lower Case Conversion Performance

- Time quadruples when double string length
- Quadratic performance







# Convert Loop To Goto Form

```
void lower(char *s)
{
    size_t i = 0;
    if (i >= strlen(s))
        goto done;
loop:
    if (s[i] >= 'A' && s[i] <= 'Z')
        s[i] -= ('A' - 'a');
    i++;
    if (i < strlen(s))
        goto loop;
done:
}
```

- strlen executed every iteration



# Calling Strlen

```
/* My version of strlen */
size_t strlen(const char *s)
{
    size_t length = 0;
    while (*s != '\0') {
        s++;
        length++;
    }
    return length;
}
```

- Strlen performance
  - Only way to determine length of string is to scan its entire length, looking for null character.
- Overall performance, string of length N
  - N calls to strlen
  - Require times N, N-1, N-2, ..., 1
  - Overall  $O(N^2)$  performance



# Improving Performance


```
void lower(char *s)
{
    size_t i;
    size_t len = strlen(s);
    for (i = 0; i < len; i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

- Move call to `strlen` outside of loop
- Since result does not change from one iteration to another
- Form of code motion



# Lower Case Conversion Performance

- Time doubles when double string length
- Linear performance of lower2





# Optimization Blocker: Procedure Calls

- *Why couldn't compiler move `strlen` out of inner loop?*
  - Procedure may have side effects
    - Alters global state each time called
  - Function may not return same value for given arguments
    - Depends on other parts of global state
    - Procedure `lower` could interact with `strlen`
- **Warning:**
  - Compiler treats procedure call as a black box
  - Weak optimizations near them
- Remedies:
  - Use of inline functions
    - GCC does this with `-O1`
      - Within single file
  - Do your own code motion

```
size_t lencnt = 0;
size_t strlen(const char *s)
{
    size_t length = 0;
    while (*s != '\0') {
        s++; length++;
    }
    lencnt += length;
    return length;
}
```

*Which one is better?*

```
long f()  
long func1() {  
    return f()+f()+f()+f();  
}  
  
long func2() {  
    return 4*f();  
}
```

A

**func1**

B

**func2**

提交



# Optimization Blocker: Procedure Calls

- Which one is better?

```
long f()  
long func1(){  
    return f()+f()+f()+f();  
}  
  
long func2(){  
    return 4*f();  
}
```

- ◆ 函数f()中改变counter值，调用次数不同函数返回值不同。
- ◆ func1()中f()被调用4次，函数返回值分别为0, 1, 2, 3, 因此func1()返回值为0+1+2+3=6;
- ◆ func2()中只调用f()一次，因此返回值为4\*0=0

- But if ...

```
long counter = 0;  
long f(){  
    return counter++;  
}
```



# Optimization Blocker: Procedure Calls

- 内联函数替换函数调用
  - 使用inline关键字定义内联函数，编译器优化时将函数调用替换为内联函数体

```
inline long f(){  
    return counter++;  
}
```

```
long func1(){  
    return f()+f()+f()+f();  
}
```

```
Long func1in(){  
    long t=counter++; //0  
    t+=counter++; //1  
    t+=counter++; //2  
    t+=counter++; //3  
    return t;  
}
```

- ◆ 通过内联函数替换掉对函数f()的四次调用
- ◆ 编译器统一func1in中对全局变量counter的更新，产生优化版本funclopt();

```
[nsrc@localhost ~]$ inlinefcall  
inline f call:6  
f call:6
```

```
Long funclopt(){  
    long t=4*counter+6; //6  
    counter+=4; //4  
    return t;  
}
```





## 练习题5.3

```
long min(long x, long y) {
    cmin++;
    return x < y ? x : y;
}

long max(long x, long y) {
    cmax++;
    return x < y ? y : x;
}

void incr(long *xp, long v) {
    cincr++;
    *xp += v;
}

long square(long x) {
    csqu++;
    return x * x;
}
```

- A.
- ```
for (i = min(x, y); i < max(x, y); incr(&i, 1))
    t += square(i);
```
- B.
- ```
for (i = max(x, y); i >= min(x, y); incr(&i, -1))
    t += square(i);
```
- C.
- ```
long low = min(x, y);
long high = max(x, y);
for (i = low; i < high; incr(&i, 1))
    t += square(i);
```

### 改进1：使用内联函数

```
inline
void inline_incr(long *xp, long v) {
    cincr++;
    *xp += v;
}

inline long inline_square(long x) {
    csqu++;
    return x * x;
}
```

### 改进2：使用宏定义

```
#define SQUA(X) (X*X)
```

Which one is better?

```
void twiddle1()  
{  
    *xp += *yp;  
    *xp += *yp;  
}  
  
void twiddle2()  
{  
    *xp += 2**yp;  
}
```

A

twiddle1

B

twiddle2

提交



# Optimization Blocker:

- Which one is better?

```
void twiddle1()  
{  
    *xp += *yp;  
    *xp += *yp;  
}  
  
void twiddle2()  
{  
    *xp += 2**yp;  
}
```

- But if  $xp = yp$ ?

```
twiddle1 result:6  
twiddle2 result:6  
case xp=yp  
twiddle1 result:8  
twiddle2 result:6
```



# Memory Matters

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        b[i] = 0;
        for (j = 0; j < n; j++)
            b[i] += a[i*n + j];
    }
}
```

```
# sum_rows1 inner loop
.L4:
    movsd    (%rsi,%rax,8), %xmm0    # FP load
    addsd    (%rdi), %xmm0           # FP add
    movsd    %xmm0, (%rsi,%rax,8)    # FP store
    addq     $8, %rdi
    cmpq     %rcx, %rdi
    jne      .L4
```

- Code updates `b[i]` on every iteration
- Why couldn't compiler optimize this away?



# Memory Aliasing

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        b[i] = 0;
        for (j = 0; j < n; j++)
            b[i] += a[i*n + j];
    }
}
```

```
double A[9] =
{ 0, 1, 2,
  4, 8, 16},
{ 32, 64, 128};

double B[3] = A+3;

sum_rows1(A, B, 3);
```

```
double A[9] =
{ 0, 1, 2,
  3, 22, 224},
{ 32, 64, 128};
```

## Value of B:

init: [4, 8, 16]

i = 0: [3, 8, 16]

i = 1: [3, 22, 16]

i = 2: [3, 22, 224]

- Code updates **b[i]** on every iteration
- Must consider possibility that these updates will affect program behavior



# Removing Aliasing

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows2(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        double val = 0;
        for (j = 0; j < n; j++)
            val += a[i*n + j];
        b[i] = val;
    }
}
```

```
# sum_rows2 inner loop
.L10:
    addsd    (%rdi), %xmm0      # FP load + add
    addq     $8, %rdi
    cmpq     %rax, %rdi
    jne      .L10
```

- No need to store intermediate results



# Optimization Blocker: Memory Aliasing

- Aliasing
  - Two different memory references specify single location
  - Easy to have happen in C
    - Since allowed to do address arithmetic
    - Direct access to storage structures
  - Get in habit of introducing local variables
    - Accumulating within loops
    - **Your way of telling compiler not to check for aliasing**




# 提纲

- 概述
- 常用的优化技术
  - 代码移动/预计算
  - 减小强度
  - 共享相同子表达式
  - 示例：冒泡排序
- 优化的阻碍
  - 函数调用
  - 内存别名
- 指令级并行
- 处理条件判断






# Modern CPU Design






为了防止你一家独大，必须得有第二供应商！



好的，老大。




码农翻身




兄弟，咱们俩合伙把  
IBM这个生意吃下来  
啊！

两位大哥，也拉兄  
弟一把！



创业翻身



**Windows**

**Linux**



**Mac**

**x86指令集**





# 中国芯的发源地

## 中国第一款正向设计的自主CPU(1999)






1999年12月 16位嵌入式微处理器核 50 MHz@0.35  $\mu\text{m}$   
2001年1月 32位/16位微处理器核 200MHz@0.25  $\mu\text{m}$



我的SW26010让太湖之光  
连续4年在全球TOP500超  
算排名第一！


这些产品都移植到了龙芯  
平台，欢迎使用！








哈哈，最终还是落到了我的手中，你的寿命也到此为止了。




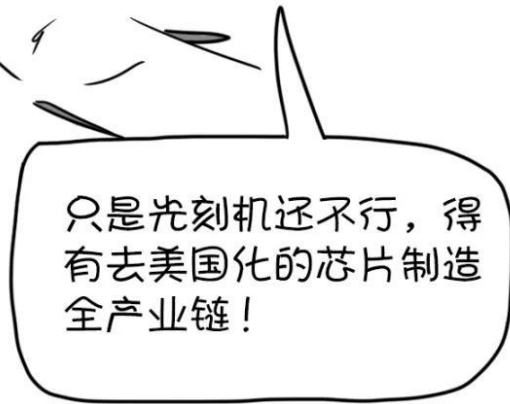


码农翻身

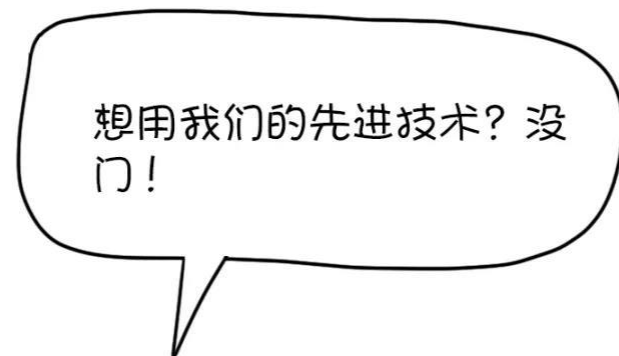

LoongArch是完全自主的指令集，可以兼容MIPS生态，融合x86/ARM指令系统的主要特点，高效支持二进制翻译。




唉，兄弟，太难了，自己的指令集，软件生态又得重来了！












# Exploiting Instruction-Level Parallelism

- Need general understanding of modern processor design
  - Hardware can execute multiple instructions in parallel
- Performance limited by data dependencies
- Simple transformations can yield dramatic performance improvement
  - Compilers often cannot make these transformations
  - Lack of associativity and distributivity in floating-point arithmetic



# Data Type for Vectors

```
/* data structure for vectors */
typedef struct{
    size_t len;
    data_t *data;
} vec;
```



## •Data Types

- Use different declarations for data\_t
- int
- long
- float
- double

```
/* retrieve vector element
   and store at val */
int get_vec_element
(*vec v, size_t idx, data_t *val)
{
    if (idx >= v->len)
        return 0;
    *val = v->data[idx];
    return 1;
}
```



# Benchmark Computation

```
void combine1(vec_ptr v, data_t *dest)
{
    long int i;
    *dest = IDENT;
    for (i = 0; i < vec_length(v); i++) {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OP val;
    }
}
```

Compute sum or  
product of vector  
elements

## •Data Types

- Use different declarations for data\_t
- int
- long
- float
- double


## •Operations

- Use different definitions of OP and IDENT
- + / 0
- \* / 1



# Cycles Per Element (CPE)

- Convenient way to express performance of program that operates on vectors or lists
- Length =  $n$
- In our case: **CPE = cycles per OP**
- $T = CPE * n + \text{Overhead}$ 
  - CPE is slope of line





# Benchmark Performance

```
void combine1(vec_ptr v, data_t *dest)
{
    long int i;
    *dest = IDENT;
    for (i = 0; i < vec_length(v); i++) {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OP val;
    }
}
```

Compute sum or  
product of vector  
elements

| Method               | Integer |       | Double FP |       |
|----------------------|---------|-------|-----------|-------|
| Operation            | Add     | Mult  | Add       | Mult  |
| Combine1 unoptimized | 22.68   | 20.02 | 19.98     | 20.18 |
| Combine1 -O1         | 10.12   | 10.12 | 10.17     | 11.14 |
| Combine1 -O3         | 4.5     | 4.5   | 6         | 7.8   |

Results in CPE (cycles per element)





# Basic Optimizations

```
void combine4(vec_ptr v, data_t *dest)
{
    long i;
    long length = vec_length(v);
    data_t *d = get_vec_start(v);
    data_t t = IDENT;
    for (i = 0; i < length; i++)
        t = t OP d[i];
    *dest = t;
}
```

- Move vec\_length out of loop
- Avoid bounds check on each cycle
- Accumulate in temporary



# Effect of Basic Optimizations

```
void combine4(vec_ptr v, data_t *dest)
{
    long i;
    long length = vec_length(v);
    data_t *d = get_vec_start(v);
    data_t t = IDENT;
    for (i = 0; i < length; i++)
        t = t OP d[i];
    *dest = t;
}
```

| Method       | Integer |       | Double FP |       |
|--------------|---------|-------|-----------|-------|
| Operation    | Add     | Mult  | Add       | Mult  |
| Combine1 -O1 | 10.12   | 10.12 | 10.17     | 11.14 |
| Combine4     | 1.27    | 3.01  | 3.01      | 5.01  |

- Eliminates sources of overhead in loop




# Superscalar Processor

- **Definition:** A superscalar processor can issue and execute *multiple instructions in one cycle*. The instructions are retrieved from a sequential instruction stream and are usually scheduled dynamically.
- **Benefit:** without programming effort, superscalar processor can take advantage of the *instruction level parallelism* that most programs have
- Most modern CPUs are superscalar.
- Intel: since Pentium (1993)



# Pipelined Functional Units

```
long mult_eg(long a, long b, long c) {  
    long p1 = a*b;  
    long p2 = a*c;  
    long p3 = p1 * p2;  
    return p3;  
}
```



| Time    |     |     |     |     |       |       |       |
|---------|-----|-----|-----|-----|-------|-------|-------|
|         | 1   | 2   | 3   | 4   | 5     | 6     | 7     |
| Stage 1 | a*b | a*c |     |     | p1*p2 |       |       |
| Stage 2 |     | a*b | a*c |     |       | p1*p2 |       |
| Stage 3 |     |     | a*b | a*c |       |       | p1*p2 |

- Divide computation into stages
- Pass partial computations from stage to stage
- Stage i can start on new computation once values passed to i+1
- E.g., complete 3 multiplications in 7 cycles, even though each requires 3 cycles



# Haswell CPU

- 8 Total Functional Units
- Multiple instructions can execute in parallel
  - 2 load, with address computation
  - 1 store, with address computation
  - 4 integer
  - 2 FP multiply
  - 1 FP add
  - 1 FP divide
- Some instructions take  $> 1$  cycle, but can be pipelined

| <i><b>Instruction</b></i>      | <i><b>Latency</b></i> | <i><b>Cycles/Issue</b></i> |
|--------------------------------|-----------------------|----------------------------|
| Load / Store                   | 4                     | 1                          |
| Integer Multiply               | 3                     | 1                          |
| <b>Integer/Long Divide</b>     | <b>3-30</b>           | <b>3-30</b>                |
| Single/Double FP Multiply      | 5                     | 1                          |
| Single/Double FP Add           | 3                     | 1                          |
| <b>Single/Double FP Divide</b> | <b>3-15</b>           | <b>3-15</b>                |



# x86-64 Compilation of Combine4

- Inner Loop (Case: Integer Multiply)

```
.L519:                                # Loop:
    imull    (%rax,%rdx,4), %ecx      # t = t * d[i]
    addq     $1, %rdx                # i++
    cmpq     %rdx, %rbp              # Compare length:i
    jg       .L519                  # If >, goto Loop
```

| Method        | Integer |      | Double FP |      |
|---------------|---------|------|-----------|------|
| Operation     | Add     | Mult | Add       | Mult |
| Combine4      | 1.27    | 3.01 | 3.01      | 5.01 |
| Latency Bound | 1.00    | 3.00 | 3.00      | 5.00 |




# Combine4 = Serial Computation (OP = \*)

- Computation (length=8)

$(((((1 * d[0]) * d[1]) * d[2]) * d[3]) * d[4]) * d[5]) * d[6]) * d[7])$

- Sequential dependence


- Performance: determined by latency of OP






# 处理器操作的抽象模型

关键路径



- 函数combine4内循环的n次迭代数据流图
  - 两条数据相关链：mul和add对程序值acc和data+i的修改
  - 假设乘法延迟为5个周期，加法延迟为1个周期
    - 左边的链成为关键路径，需要5n个周期执行
    - 右边的链需要n个周期执行，不会制约程序性能
- 浮点乘法器成为程序制约资源，其他操作与乘法器并行执行







# 循环展开Loop Unrolling (2x1)

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = (x OP d[i]) OP d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OP d[i];
    }
    *dest = x;
}
```

- Perform 2x more useful work per iteration



# Effect of Loop Unrolling

| Method               | Integer     |             | Double FP   |             |
|----------------------|-------------|-------------|-------------|-------------|
| Operation            | Add         | Mult        | Add         | Mult        |
| Combine4             | 1.27        | 3.01        | 3.01        | 5.01        |
| Unroll 2x1           | 1.01        | 3.01        | 3.01        | 5.01        |
| <i>Latency Bound</i> | <i>1.00</i> | <i>3.00</i> | <i>3.00</i> | <i>5.00</i> |

- Helps integer add
  - Achieves latency bound
- Others don't improve. *Why?*
  - Still sequential dependency

```
x = (x OP d[i]) OP d[i+1];
```



# Loop Unrolling with Reassociation (2x1a)

```
void unroll2aa_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = x OP (d[i] OP d[i+1]);
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OP d[i];
    }
    *dest = x;
}
```

Compare to before


$x = (x \text{ OP } d[i]) \text{ OP } d[i+1];$

- Can this change the result of the computation?
- Yes, for FP. *Why?*



# Reassociated Computation

```
x = x OP (d[i] OP d[i+1]);
```



- What changed:
  - Ops in the next iteration can be started early (no dependency)
- Overall Performance
  - N elements, D cycles latency/op
  - $(N/2+1)*D$  cycles:  
 **$CPE = D/2$**



# Effect of Reassociation

| Method                  | Integer     |             | Double FP   |             |
|-------------------------|-------------|-------------|-------------|-------------|
| Operation               | Add         | Mult        | Add         | Mult        |
| Combine4                | 1.27        | 3.01        | 3.01        | 5.01        |
| Unroll 2x1              | 1.01        | 3.01        | 3.01        | 5.01        |
| Unroll 2x1a             | 1.01        | 1.51        | 1.51        | 2.51        |
| <i>Latency Bound</i>    | <i>1.00</i> | <i>3.00</i> | <i>3.00</i> | <i>5.00</i> |
| <i>Throughput Bound</i> | <i>0.50</i> | <i>1.00</i> | <i>1.00</i> | <i>0.50</i> |

4 func. units for int +,  
2 func. units for load

2 func. units for FP \*,  
2 func. units for load

- Nearly 2x speedup for Int \*, FP +, FP \*
- Reason: Breaks sequential dependency

```
x = x OP (d[i] OP d[i+1]);
```



# Loop Unrolling with Separate Accumulators (2x2)

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x0 = IDENT;
    data_t x1 = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x0 = x0 OP d[i];
        x1 = x1 OP d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x0 = x0 OP d[i];
    }
    *dest = x0 OP x1;
}
```

- Different form of reassociation



# Effect of Separate Accumulators

| Method                  | Integer     |             | Double FP   |             |
|-------------------------|-------------|-------------|-------------|-------------|
| Operation               | Add         | Mult        | Add         | Mult        |
| Combine4                | 1.27        | 3.01        | 3.01        | 5.01        |
| Unroll 2x1              | 1.01        | 3.01        | 3.01        | 5.01        |
| Unroll 2x1a             | 1.01        | 1.51        | 1.51        | 2.51        |
| Unroll 2x2              | 0.81        | 1.51        | 1.51        | 2.51        |
| <i>Latency Bound</i>    | <i>1.00</i> | <i>3.00</i> | <i>3.00</i> | <i>5.00</i> |
| <i>Throughput Bound</i> | <i>0.50</i> | <i>1.00</i> | <i>1.00</i> | <i>0.50</i> |

- Int + makes use of two load units

```
x0 = x0 OP d[i];  
x1 = x1 OP d[i+1];
```

- 2x speedup (over unroll2) for Int \*, FP +, FP \*



Diagram illustrating a binary tree structure for a sequence of data points  $d_0, d_1, \dots, d_7$ . The tree consists of 8 internal nodes (represented by pink boxes with an asterisk  $*$ ) and 9 leaf nodes (represented by downward arrows). The root node takes inputs 1 and  $d_0$ . The tree structure is as follows:

- Root node (Level 0) takes inputs 1 and  $d_0$ .
- Level 1: The root node's output goes to the left child, which takes input  $d_2$ . The root node's output also goes to the right child, which takes inputs 1 and  $d_1$ .
- Level 2: The left child of the root node's left child takes input  $d_4$ . The right child of the root node's left child takes input  $d_3$ .
- Level 3: The left child of the left child of the root node's left child takes input  $d_6$ . The right child of the left child of the root node's left child takes input  $d_5$ .
- Level 4: The left child of the left child of the left child of the root node's left child takes input  $d_7$ .
- The final output is from the bottom-most node.

- Two independent “streams” of operations

- N elements, D cycles latency/op
- Should be  $(N/2+1)*D$  cycles:  
**CPE = D/2**
- CPE matches prediction!

## What Now?





# Unrolling & Accumulating

- Idea
  - Can unroll to any degree  $L$
  - Can accumulate  $K$  results in parallel
  - $L$  must be multiple of  $K$
- Limitations
  - Diminishing returns
    - Cannot go beyond throughput limitations of execution units
  - Large overhead for short lengths
    - Finish off iterations sequentially



# Unrolling & Accumulating: Double \*

- Case
  - Intel Haswell
  - Double FP Multiplication
  - Latency bound: 5.00. Throughput bound: 0.50

Accumulators

| FP * | Unrolling Factor L |      |      |      |      |      |      |      |
|------|--------------------|------|------|------|------|------|------|------|
| K    | 1                  | 2    | 3    | 4    | 6    | 8    | 10   | 12   |
| 1    | 5.01               | 5.01 | 5.01 | 5.01 | 5.01 | 5.01 | 5.01 |      |
| 2    |                    | 2.51 |      | 2.51 |      | 2.51 |      |      |
| 3    |                    |      | 1.67 |      |      |      |      |      |
| 4    |                    |      |      | 1.25 |      | 1.26 |      |      |
| 6    |                    |      |      |      | 0.84 |      |      | 0.88 |
| 8    |                    |      |      |      |      | 0.63 |      |      |
| 10   |                    |      |      |      |      |      | 0.51 |      |
| 12   |                    |      |      |      |      |      |      | 0.52 |



# Unrolling & Accumulating: Int +

- Case
  - Intel Haswell
  - Integer addition
  - Latency bound: 1.00. Throughput bound: 1.00

Accumulators

| Int + | Unrolling Factor L |      |      |      |      |      |      |      |
|-------|--------------------|------|------|------|------|------|------|------|
| K     | 1                  | 2    | 3    | 4    | 6    | 8    | 10   | 12   |
| 1     | 1.27               | 1.01 | 1.01 | 1.01 | 1.01 | 1.01 | 1.01 |      |
| 2     |                    | 0.81 |      | 0.69 |      | 0.54 |      |      |
| 3     |                    |      | 0.74 |      |      |      |      |      |
| 4     |                    |      |      | 0.69 |      | 1.24 |      |      |
| 6     |                    |      |      |      | 0.56 |      |      | 0.56 |
| 8     |                    |      |      |      |      | 0.54 |      |      |
| 10    |                    |      |      |      |      |      | 0.54 |      |
| 12    |                    |      |      |      |      |      |      | 0.56 |



# Achievable Performance

| Method                  | Integer     |             | Double FP   |             |
|-------------------------|-------------|-------------|-------------|-------------|
| Operation               | Add         | Mult        | Add         | Mult        |
| Best                    | 0.54        | 1.01        | 1.01        | 0.52        |
| <i>Latency Bound</i>    | <i>1.00</i> | <i>3.00</i> | <i>3.00</i> | <i>5.00</i> |
| <i>Throughput Bound</i> | <i>0.50</i> | <i>1.00</i> | <i>1.00</i> | <i>0.50</i> |

- Limited only by throughput of functional units
- Up to 42X improvement over original, unoptimized code



# Programming with AVX2

## YMM Registers

■ 16 total, each 32 bytes

■ 32 single-byte integers



■ 16 16-bit integers



■ 8 32-bit integers



■ 8 single-precision floats



■ 4 double-precision floats



■ 1 single-precision float



■ 1 double-precision float






# 单指令多数据


## ■ SIMD Operations: Single Precision

**vaddsd %ymm0, %ymm1, %ymm1**



## ■ SIMD Operations: Double Precision

**vaddpd %ymm0, %ymm1, %ymm1**





# 向量指令

| Method                      | Integer     |             | Double FP   |             |
|-----------------------------|-------------|-------------|-------------|-------------|
| Operation                   | Add         | Mult        | Add         | Mult        |
| Scalar Best                 | 0.54        | 1.01        | 1.01        | 0.52        |
| Vector Best                 | 0.06        | 0.24        | 0.25        | 0.16        |
| <i>Latency Bound</i>        | <i>0.50</i> | <i>3.00</i> | <i>3.00</i> | <i>5.00</i> |
| <i>Throughput Bound</i>     | <i>0.50</i> | <i>1.00</i> | <i>1.00</i> | <i>0.50</i> |
| <i>Vec Throughput Bound</i> | <i>0.06</i> | <i>0.12</i> | <i>0.25</i> | <i>0.12</i> |

- Make use of AVX Instructions
  - Parallel operations on multiple data elements
  - See Web Aside OPT:SIMD on CS:APP web page



# What About Branches?

- Challenge
  - **Instruction Control Unit** must work well ahead of **Execution Unit** to generate enough operations to keep EU busy

```
404663:  mov    $0x0,%eax
404668:  cmp    (%rdi),%rsi
40466b:  jge    404685
40466d:  mov    0x8(%rdi),%rax

. . .

404685:  repz   retq
```

} Executing


← How to continue?

- When encounters conditional branch, cannot reliably determine where to continue fetching





# Modern CPU Design





# Branch Outcomes

- **When encounter conditional branch, cannot determine where to continue fetching**
  - Branch Taken: Transfer control to branch target
  - Branch Not-Taken: Continue with next instruction in sequence
- **Cannot resolve until outcome determined by branch/integer unit**

```
404663:  mov    $0x0,%eax
404668:  cmp    (%rdi),%rsi
40466b:  jge    404685
40466d:  mov    0x8(%rdi),%rax

. . .

404685:  repz   retq
```

Branch Not-Taken

Branch Taken



# Branch Prediction

- Idea
  - Guess which way branch will go
  - Begin executing instructions at predicted position
    - But don't actually modify register or memory data

```
404663:  mov    $0x0,%eax
404668:  cmp    (%rdi),%rsi
40466b:  jge    404685
40466d:  mov    0x8(%rdi),%rax

. . .

404685:  repz   retq
```

**Predict Taken**

} **Begin  
Execution**



# Branch Prediction Through Loop

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029
```

*i = 98*

Assume  
vector length = 100

Predict Taken (OK)

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029
```

*i = 99*

Predict Taken  
(Oops)


```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029
```

*i = 100*

Read  
invalid  
location

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029
```

*i = 101*





# Branch Misprediction Invalidation

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029 i = 98
```

Assume  
vector length = **100**

Predict Taken (OK)

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029 i = 99
```

Predict Taken  
(Oops)

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029 i = 100
```

Invalidate

```
401029: vmulsd (%rdx), %xmm0, %xmm0
40102d: add    $0x8, %rdx
401031: cmp    %rax, %rdx
401034: jne    401029 i = 101
```



# Branch Misprediction Recovery

```
401029: vmulsd (%rdx), %xmm0, %xmm0
```

```
40102d: add     $0x8, %rdx
```

```
401031: cmp     %rax, %rdx
```

```
401034: jne     401029
```

```
401036: jmp     401040
```

```
. . .
```

```
401040: vmovsd %xmm0, (%r12)
```

*i = 99*

Definitely not taken

Reload  
Pipeline

- Performance Cost
  - Multiple clock cycles on modern processor
  - Can be a major performance limiter



# Branch Prediction Numbers

- Default behavior:
  - Backwards branches are often loops so predict taken
  - Forwards branches are often if so predict not taken
- Predictors average better than 95% accuracy
  - Most branches are already predictable.
- Bonus material:  
<http://stackoverflow.com/questions/11227809/why-is-processing-a-sorted-array-faster-than-an-unsorted-array>



# 确信和消除性能瓶颈

## 程序剖析

- 程序剖析(profiling): 是程序执行时收集性能数据的分析工具, 方法是运行程序的一个版本, 其中插入了工具代码, 以确定程序的各个部分需要多少时间

- GRPOF: 给出函数执行时间和调用次数

1. 为剖析编译和链接程序: 使用-pg选项

- 例: `gcc -O3 -pg prefixsum.c -o prefixsum`

2. 执行程序

- 例: `prefixsum`

3. 调用GPROF分析

- 例: `gprof prefixsum`,  
剖析报告显示函数执行信息

granularity: each sample hit covers 2 byte(s) for 0.15% of 6.86 seconds

| index         | % time | self | children | called | name       |
|---------------|--------|------|----------|--------|------------|
| <spontaneous> |        |      |          |        |            |
| [1]           | 100.0  | 1.38 | 5.48     |        | main [1]   |
|               |        | 2.87 | 0.00     | 1/1    | psuml [2]  |
|               |        | 1.70 | 0.00     | 1/1    | psum2 [3]  |
|               |        | 0.91 | 0.00     | 1/1    | psumla [4] |
| -----         |        |      |          |        |            |
|               |        | 2.87 | 0.00     | 1/1    | main [1]   |
| [2]           | 41.9   | 2.87 | 0.00     | 1      | psuml [2]  |
| -----         |        |      |          |        |            |
|               |        | 1.70 | 0.00     | 1/1    | main [1]   |
| [3]           | 24.8   | 1.70 | 0.00     | 1      | psum2 [3]  |
| -----         |        |      |          |        |            |
|               |        | 0.91 | 0.00     | 1/1    | main [1]   |
| [4]           | 13.2   | 0.91 | 0.00     | 1      | psumla [4] |
| -----         |        |      |          |        |            |





# perf监测性能指标

- Perf 是用来进行软件性能分析的工具，可以同时分析应用代码和内核，从而全面理解应用程序中的性能瓶颈。
- Perf程序剖析示例：监测rearrange程序branch与cache misses
  - 命令：perf stat -e branches -e branch-misses -e cache-misses ./rearrange

```
[nsrc@localhost ~]$ perf stat -e branches -e branch-misses -e cache-misses ./rearrange
Input length of random array a and b:1000000
Input filter ratio of array:1
numbers of 1:1000000
minmax1 cycles_per_elem:17.6
minmax1 cycles_per_elem: 2.7
bredicate branching: Items:1000000, cycles_per_elem: 5.9
bredicate branching: Items:1000000, cycles_per_elem: 3.9

Performance counter stats for './rearrange':

      56,337,543      branches
      613,754        branch-misses          #    1.09% of all branches
         3,327        cache-misses

      3.949698873 seconds time elapsed
```

- Perf主要监测指标：perf list

```
[nsrc@localhost ~]$ perf list
List of pre-defined events (to be used in -e):

branch-instructions OR branches      [Hardware event]
branch-misses                        [Hardware event]
bus-cycles                          [Hardware event]
cache-misses                        [Hardware event]
cache-references                    [Hardware event]
cpu-cycles OR cycles                [Hardware event]
instructions                        [Hardware event]
ref-cycles                          [Hardware event]

alignment-faults                    [Software event]
context-switches OR cs              [Software event]
cpu-clock                           [Software event]
cpu-migrations OR migrations        [Software event]
dummy                               [Software event]
emulation-faults                    [Software event]
major-faults                        [Software event]
minor-faults                        [Software event]
page-faults OR faults               [Software event]
task-clock                          [Software event]

L1-dcache-load-misses               [Hardware cache event]
L1-dcache-loads                     [Hardware cache event]
L1-dcache-stores                    [Hardware cache event]
L1-icache-load-misses               [Hardware cache event]
LLC-loads                           [Hardware cache event]
LLC-load-misses                     [Hardware cache event]
LLC-stores                          [Hardware cache event]
branch-load-misses                  [Hardware cache event]
branch-loads                        [Hardware cache event]
dTLB-load-misses                    [Hardware cache event]
dTLB-loads                          [Hardware cache event]
dTLB-store-misses                   [Hardware cache event]
dTLB-stores                         [Hardware cache event]
iTLB-load-misses                    [Hardware cache event]
iTLB-loads                          [Hardware cache event]
node-load-misses                    [Hardware cache event]
node-loads                          [Hardware cache event]
node-store-misses                   [Hardware cache event]
node-stores                         [Hardware cache event]

branch-instructions OR cpu/branch-instructions/ [Kernel PMU event]
branch-misses OR cpu/branch-misses/           [Kernel PMU event]
bus-cycles OR cpu/bus-cycles/                  [Kernel PMU event]
cache-misses OR cpu/cache-misses/              [Kernel PMU event]
cache-references OR cpu/cache-references/       [Kernel PMU event]
cpu-cycles OR cpu/cpu-cycles/                  [Kernel PMU event]
cycles-ct OR cpu/cycles-ct/                    [Kernel PMU event]
cycles-t OR cpu/cycles-t/                      [Kernel PMU event]
```



# perf监测性能指标

- Perf性能分析：监测程序性能并将其数据保存到perf.data中，使用perf report进行分析。
  - perf record和perf report可以更精确的分析一个应用，perf record可以精确到函数级别。并且在函数里面混合显示汇编语言和代码。
- 命令：
  - perf record -e branches -e branch-misses -e cache-misses ./rearrange

|                                                                    |           |                   |                         |
|--------------------------------------------------------------------|-----------|-------------------|-------------------------|
| 202.112.117.222 - Xshell 4                                         |           |                   |                         |
| 文件(F) 编辑(E) 查看(V) 工具(T) 窗口(W) 帮助(H)                                |           |                   |                         |
| 新建 重新连接                                                            |           |                   |                         |
| 1 202.112.117.222                                                  |           |                   |                         |
| Samples: 301 of event 'branches', Event count (approx.): 7449      |           |                   |                         |
| Overhead                                                           | Command   | Shared Object     | Symbol                  |
| 32.91%                                                             | rearrange | rearrange         | [.] minmax1             |
| 19.59%                                                             | rearrange | libc-2.17.so      | [.] main                |
| 18.09%                                                             | rearrange | libc-2.17.so      | [.] rand                |
| 10.08%                                                             | rearrange | libc-2.17.so      | [.] __random            |
| 5.30%                                                              | rearrange | rearrange         | [.] __dl_addr           |
| 4.39%                                                              | rearrange | rearrange         | [k] rebalance_domains   |
| 3.34%                                                              | rearrange | rearrange         | [k] strnlen_user        |
| 1.72%                                                              | rearrange | rearrange         | [k] prepend_path        |
| 1.66%                                                              | rearrange | rearrange         | [.] __random_r          |
| 0.66%                                                              | rearrange | [kernel.kallsyms] | [k] get_page_from_freel |
| 0.38%                                                              | rearrange | [kernel.kallsyms] | [k] ktime_get_update_of |
| 0.35%                                                              | rearrange | [kernel.kallsyms] | [k] find_yma            |
| 0.30%                                                              | rearrange | [kernel.kallsyms] | [k] smp_call_function_s |
| 0.28%                                                              | rearrange | [kernel.kallsyms] | [k] native_apic_mem_wri |
| 0.27%                                                              | rearrange | libc-2.17.so      | [k] update_cfs_rq_block |
| 0.26%                                                              | rearrange | [kernel.kallsyms] | [k] perf_event_comm_out |
| 0.09%                                                              | rearrange | [kernel.kallsyms] | [k] perf_event_comm_out |
| 0.09%                                                              | rearrange | [kernel.kallsyms] | [k] perf_pmu_enable     |
| 0.08%                                                              | rearrange | [kernel.kallsyms] | [k] perf_pmu_enable     |
| 0.08%                                                              | rearrange | [kernel.kallsyms] | [k] perf_pmu_enable     |
| 0.00%                                                              | perf      | [kernel.kallsyms] | [k] perf_pmu_enable     |
| p: If you prefer Intel style assembly, try: perf annotate -M intel |           |                   |                         |
| 已连接 202.112.117.222:2017. SSH2 xterm 72x25 20,72 1 会话              |           |                   |                         |

|                                                                    |           |                   |                         |
|--------------------------------------------------------------------|-----------|-------------------|-------------------------|
| 202.112.117.222 - Xshell 4                                         |           |                   |                         |
| 文件(F) 编辑(E) 查看(V) 工具(T) 窗口(W) 帮助(H)                                |           |                   |                         |
| 新建 重新连接                                                            |           |                   |                         |
| 1 202.112.117.222                                                  |           |                   |                         |
| Samples: 265 of event 'branch-misses', Event count (approx.): 7449 |           |                   |                         |
| Overhead                                                           | Command   | Shared Object     | Symbol                  |
| 82.36%                                                             | rearrange | rearrange         | [.] minmax1             |
| 14.23%                                                             | rearrange | rearrange         | [.] main                |
| 1.07%                                                              | rearrange | libc-2.17.so      | [.] rand                |
| 0.84%                                                              | rearrange | libc-2.17.so      | [.] __random            |
| 0.38%                                                              | rearrange | libc-2.17.so      | [.] __dl_addr           |
| 0.29%                                                              | rearrange | [kernel.kallsyms] | [k] rebalance_domains   |
| 0.18%                                                              | rearrange | [kernel.kallsyms] | [k] strnlen_user        |
| 0.11%                                                              | rearrange | [kernel.kallsyms] | [k] prepend_path        |
| 0.09%                                                              | rearrange | libc-2.17.so      | [.] __random_r          |
| 0.08%                                                              | rearrange | [kernel.kallsyms] | [k] get_page_from_freel |
| 0.07%                                                              | rearrange | [kernel.kallsyms] | [k] ktime_get_update_of |
| 0.07%                                                              | rearrange | [kernel.kallsyms] | [k] find_yma            |
| 0.07%                                                              | rearrange | [kernel.kallsyms] | [k] smp_call_function_s |
| 0.07%                                                              | rearrange | [kernel.kallsyms] | [k] native_apic_mem_wri |
| 0.06%                                                              | rearrange | [kernel.kallsyms] | [k] update_cfs_rq_block |
| 0.04%                                                              | perf      | [kernel.kallsyms] | [k] perf_event_comm_out |
| 0.01%                                                              | perf      | [kernel.kallsyms] | [k] perf_pmu_enable     |
| p: If you prefer Intel style assembly, try: perf annotate -M intel |           |                   |                         |
| 已连接 202.112.117.222:2017. SSH2 xterm 72x25 20,72 1 会话              |           |                   |                         |

|                                                                    |           |                   |                               |
|--------------------------------------------------------------------|-----------|-------------------|-------------------------------|
| 202.112.117.222 - Xshell 4                                         |           |                   |                               |
| 文件(F) 编辑(E) 查看(V) 工具(T) 窗口(W) 帮助(H)                                |           |                   |                               |
| 新建 重新连接                                                            |           |                   |                               |
| 1 202.112.117.222                                                  |           |                   |                               |
| Samples: 27 of event 'cache-misses', Event count (approx.): 7449   |           |                   |                               |
| Overhead                                                           | Command   | Shared Object     | Symbol                        |
| 45.50%                                                             | rearrange | [kernel.kallsyms] | [k] get_page_from_freelist    |
| 15.20%                                                             | rearrange | libc-2.17.so      | [.] __dl_addr                 |
| 12.73%                                                             | rearrange | [kernel.kallsyms] | [k] __mod_zone_page_state     |
| 9.67%                                                              | rearrange | [kernel.kallsyms] | [k] clear_page_c_e            |
| 5.49%                                                              | rearrange | [kernel.kallsyms] | [k] __rmqueue                 |
| 2.94%                                                              | rearrange | [kernel.kallsyms] | [k] page_fault                |
| 2.77%                                                              | rearrange | [kernel.kallsyms] | [k] copy_page_rep             |
| 2.77%                                                              | rearrange | [kernel.kallsyms] | [k] get_pageblock_flags_group |
| 1.18%                                                              | rearrange | [kernel.kallsyms] | [k] __list_del_entry          |
| 0.43%                                                              | rearrange | rearrange         | [.] minmax1                   |
| 0.40%                                                              | rearrange | libc-2.17.so      | [.] __mpn_extract_double      |
| 0.40%                                                              | rearrange | libc-2.17.so      | [.] __printf_fp               |
| 0.32%                                                              | rearrange | [kernel.kallsyms] | [k] padzero                   |
| 0.13%                                                              | rearrange | [kernel.kallsyms] | [k] perf_event_aux_ctx        |
| 0.05%                                                              | perf      | [kernel.kallsyms] | [k] perf_event_aux.part.48    |
| 0.03%                                                              | perf      | [kernel.kallsyms] | [k] perf_pmu_enable           |
| p: If you prefer Intel style assembly, try: perf annotate -M intel |           |                   |                               |
| 已连接 202.112.117.222:2017. SSH2 xterm 72x25 19,72 1 会话              |           |                   |                               |



# 代码阅读工具

Source Insight

hash join Project - Source Insight 4.0 - [main.c (C:\documents\...\src)]

File Edit Search Project Options Tools View Window Help

main.c (C:\documents\...\src) ×

**main.c**

Symbol Name (Alt+L)

- ifndef GNU\_SOURCE
- GNU\_SOURCE
- endif
- include "cpu\_mapping.h" /\* get\_cpu\_id
- include <sched.h> /\* sched\_setaffinity
- include <stdio.h> /\* printf \*/
- include <sys/time.h> /\* gettimeofday
- include <getopt.h> /\* getopt \*/
- include <stdlib.h> /\* exit \*/
- include <string.h> /\* strcmp \*/
- include <limits.h> /\* INT\_MAX \*/
- include <sys/time.h> /\* gettimeofday
- include "rdtsc.h" /\* startTimer, stopTim
- include "no\_partitioning\_join.h" /\* no p
- include "parallel\_radix\_join.h" /\* paralle
- include "generator.h" /\* create\_relatior
- include "prj\_params.h" /\* RELATION\_P
- include "perf\_counters.h" /\* PCM\_x \*/
- include "affinity.h" /\* pthread\_attr\_seta
- include "../config.h" /\* autoconf heade
- if !defined(\_\_cplusplus)
- getopt
- endif
- RAND\_RANGE
- MALLOC
- algo t
- param t
- num\_part
- num\_supplier
- num\_customer
- num\_date
- num\_lineorder
- FactColumns
- DimVector
- MeasureIndex
- firstfilterflag
- M1

\* with --perfconf which specifies which hardware counters to monitor. For  
\* detailed list of hardware counters consult to "Intel 64 and IA-32  
\* Architectures Software Developer's Manual" Appendix A. For an example  
\* configuration file used in the experiments, see <b>pcm.cfg</b> file.  
\* Lastly, an output file name with --perfout on commandline can be specified to  
\* print out profiling results, otherwise it defaults to stdout.  
\*  
\* @subsection systemparams System and Implementation Parameters  
\*  
\* The join implementations need to know about the system at hand to a certain  
\* degree. For instance #CACHE\_LINE\_SIZE is required by both of the  
\* implementations. In case of no partitioning join, other implementation  
\* parameters such as bucket size or whether to pre-allocate for overflowing  
\* buckets are parametrized and can be modified in 'npj\_params.h'.  
\*  
\* On the other hand, radix joins are more sensitive to system parameters and  
\* the optimal setting of parameters should be found from machine to machine to  
\* get the same results as presented in the paper. System parameters needed are  
\* #CACHE\_LINE\_SIZE, #L1\_CACHE\_SIZE and  
\* #L1\_ASSOCIATIVITY. Other implementation parameters specific to radix  
\* join are also important such as #NUM\_RADIX\_BITS  
\* which determines number of created partitions and #NUM\_PASSES which  
\* determines number of partitioning passes. Our implementations support between  
\* 1 and 2 passes and they can be configured using these parameters to find the  
\* ideal performance on a given machine.  
\*  
\* @section data Generating Data Sets of Our Experiments  
\*  
\* Here we briefly describe how to generate data sets used in our experiments  
\* with the command line parameters above.  
\*  
\* @subsection workloadB Workload B  
\*  
\* In this data set, the inner relation R and outer relation S have 128\*10^6  
\* tuples each. The tuples are 8 bytes long, consisting of 4-byte (or 32-bit)  
\* integers and a 4-byte payload. As for the data distribution, if not  
\* explicitly specified, we use relations with randomly shuffled unique keys  
\* ranging from 1 to 128\*10^6. To generate this data set, append the following  
\* parameters to the executable  
\* <tt>mchashjoins</tt>:  
\*  
\* @verbatim  
\* \$ ./mchashjoins [other options] --r-size=128000000 --s-size=128000000  
\* @endverbatim  
\*  
\* \note Configure must have run without --enable-key88.  
\*  
\* @subsection workloadA Workload A  
\*  
\* This data set reflects the case where the join is performed between the

Project Files × Project Symbols ×

File Name (Ctrl+O)

| File Name             | Directory     | Size |
|-----------------------|---------------|------|
| affinity.h            | C:\documents\ | 1    |
| barrier.h             | C:\documents\ | 2    |
| cpu_mapping.c         | C:\documents\ | 1    |
| cpu_mapping.h         | C:\documents\ | 1    |
| generator.c           | C:\documents\ | 19   |
| generator.h           | C:\documents\ | 3    |
| genzipf.c             | C:\documents\ | 3    |
| genzipf.h             | C:\documents\ | 1    |
| lock.h                | C:\documents\ | 1    |
| main.c                | C:\documents\ | 37   |
| Makefile              | C:\documents\ | 28   |
| mapdql_history.txt    | C:\documents\ | 10   |
| no_partitioning_join. | C:\documents\ | 61   |
| no_partitioning_join. | C:\documents\ | 2    |
| npj_params.h          | C:\documents\ | 1    |
| npj_types.h           | C:\documents\ | 53   |
| parallel_radix_join.c | C:\documents\ | 2    |
| parallel_radix_join.h | C:\documents\ | 9    |
| perf_counters.c       | C:\documents\ | 2    |
| perf_counters.h       | C:\documents\ | 2    |
| prj_params.h          | C:\documents\ | 1    |
| rdtsc.h               | C:\documents\ | 6    |
| task_queue.h          | C:\documents\ | 1    |
| types.h               | C:\documents\ | 1    |

Context

Line 175 Col 3 [UTF-8]

INS



# Getting High Performance

- Good compiler and flags
- Don't do anything stupid
  - Watch out for hidden algorithmic inefficiencies
  - Write compiler-friendly code
    - Watch out for optimization blockers:  
procedure calls & memory references
  - Look carefully at innermost loops (where most work is done)
- Tune code for machine
  - Exploit instruction-level parallelism
  - Avoid unpredictable branches
  - Make code cache friendly (Covered later in course)



预祝同学们：  
考得都会  
蒙的都对  
新年快乐

相约**2027** 《人工智能计算系统》